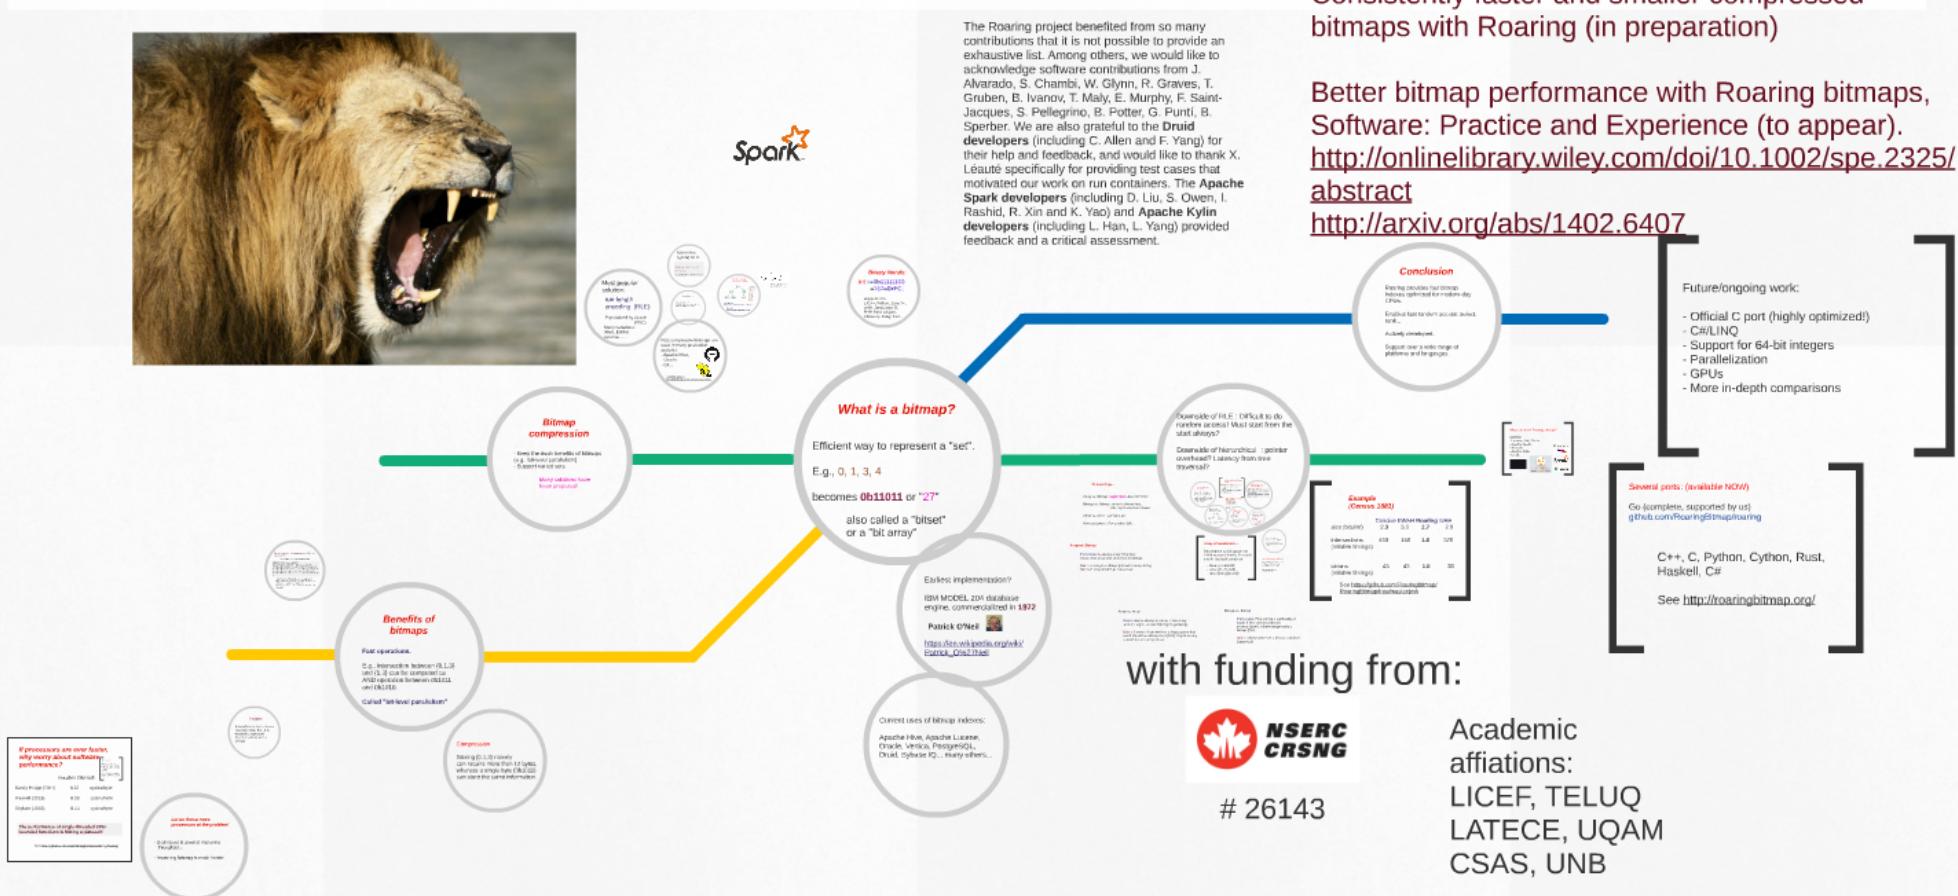


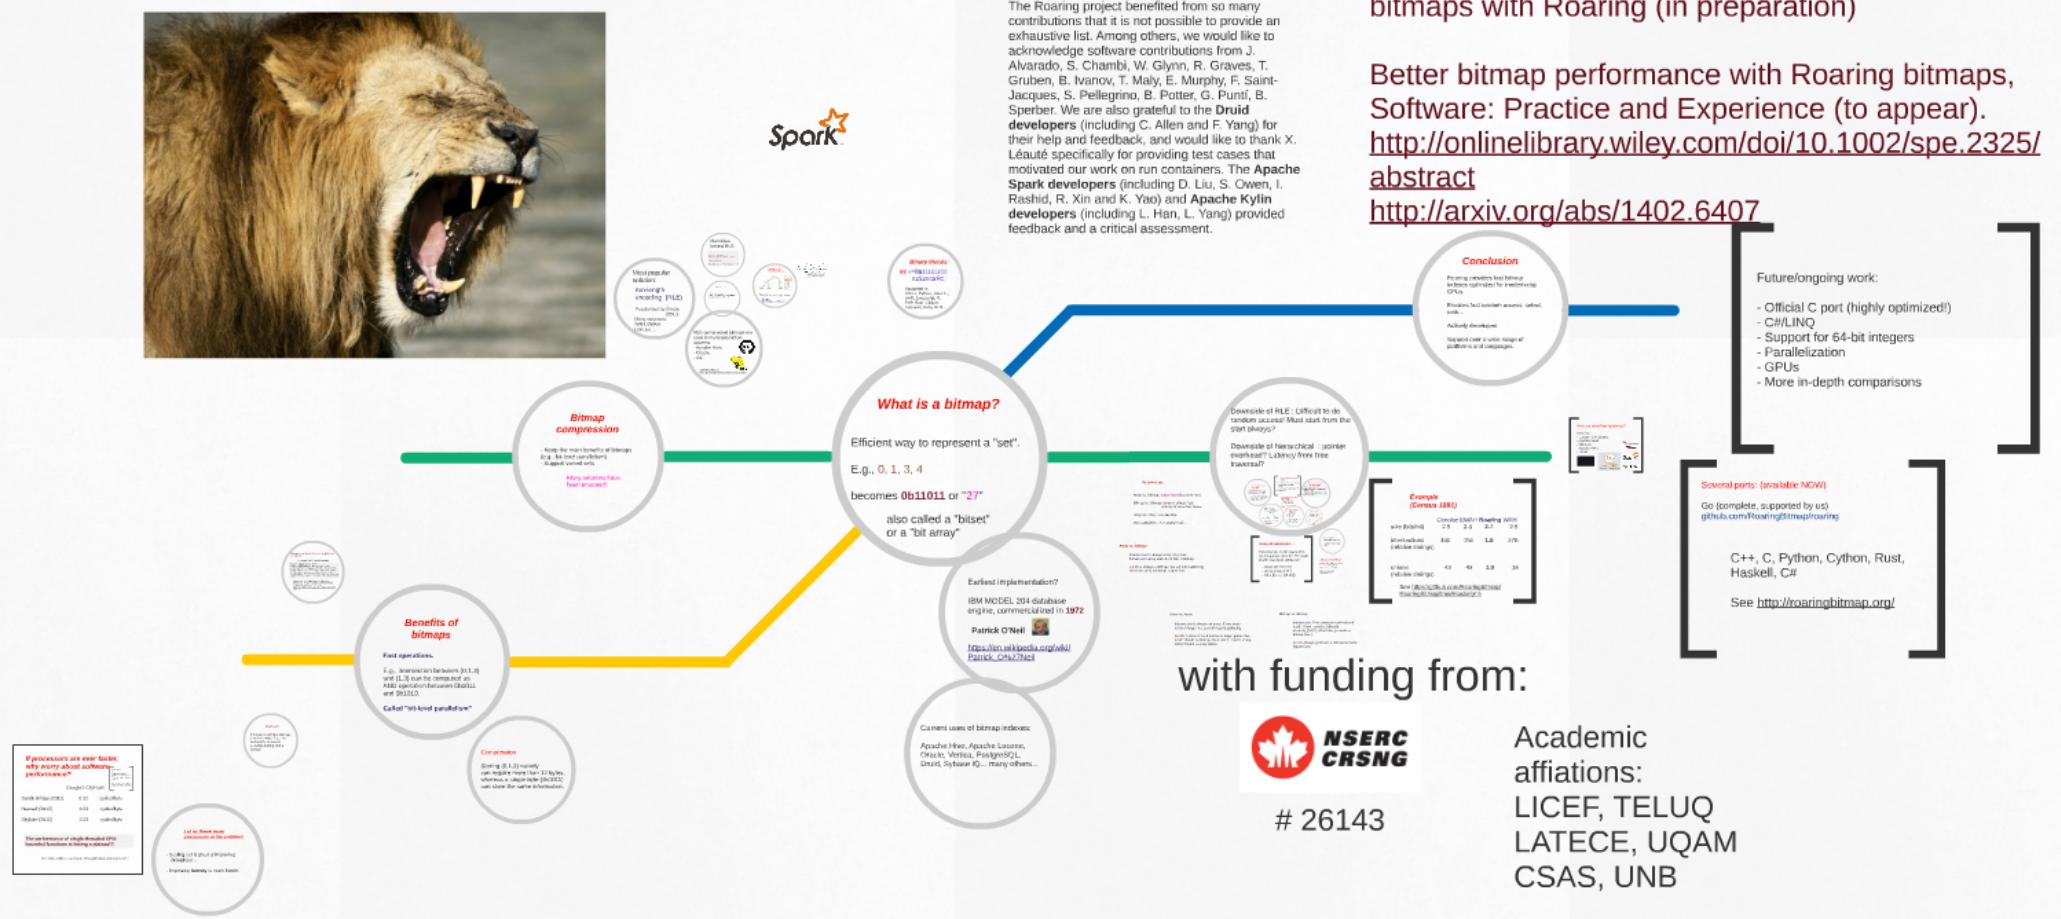
January 2016

# Roaring Bitmaps



<http://roaringbitmap.org/>  
<https://github.com/RoaringBitmap>

January 2016



<http://roaringbitmap.org/>  
<https://github.com/RoaringBitmap>

# Roaring Bitmaps

Consistently faster and smaller compressed bitmaps with Roaring (in preparation)

Better bitmap performance with Roaring bitmaps, Software: Practice and Experience (to appear).

<http://onlinelibrary.wiley.com/doi/10.1002/spe.2325/abstract>

<http://arxiv.org/abs/1402.6407>

with funding from:



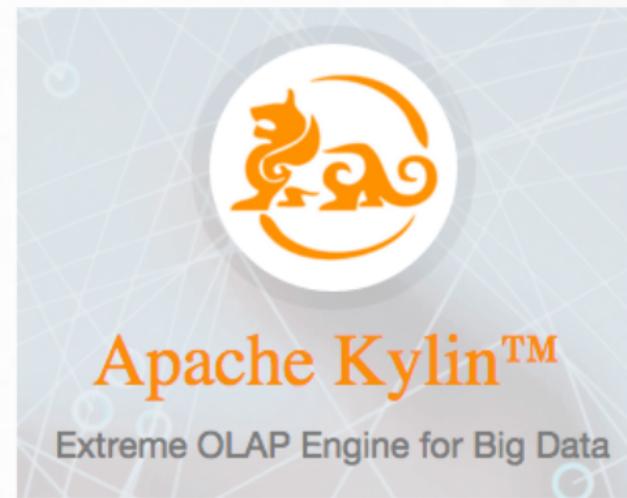
# 26143

Academic affiliations:  
LICEF, TELUQ  
LATECE, UQAM  
CSAS, UNB

# Why care about Roaring bitmaps?

Used by:

- Lucene, Solr, Elastic
- Apache Spark,
- Whoosh,
- Apache Kylin,
- Druid.



# **If processors are ever faster, why worry about software performance?**

Google's CityHash

This means:

"High performance"  
programming will remain  
very important in coming  
decade.

We have to tune software  
for the hardware we have...

Sandy Bridge (2011)	0.32	cycles/byte
Haswell (2013)	0.23	cycles/byte
Skylake (2015)	0.23	cycles/byte

**The performance of single-threaded CPU-  
bounded functions is hitting a plateau!!!**

Ref: <https://github.com/lemire/StronglyUniversalStringHashing>

## ***Binary literals:***

```
int x=0b11111100  
      =252=0xFC;
```

Supported in  
C/C++, Python, Java 7+,  
Swift, JavaScript, D,  
PHP, Rust, Clojure,  
C#(soon), Ruby, Perl...

## ***What is a bitmap?***

Efficient way to represent a "set".

E.g., 0, 1, 3, 4

becomes **0b11011** or "27"

also called a "bitset"  
or a "bit array"

Earliest implementation?

IBM MODEL 204 database  
engine, commercialized in **1972**

Patrick O'Neil



[https://en.wikipedia.org/wiki/  
Patrick\\_O%27Neil](https://en.wikipedia.org/wiki/Patrick_O%27Neil)

## ***Benefits of bitmaps***

### **Fast operations.**

E.g., intersection between {0,1,3} and {1,3} can be computed as AND operation between 0b1011 and 0b1010.

**Called "bit-level parallelism"**

## Compression

Storing  $\{0,1,3\}$  naively can require more than 12 bytes, whereas a single byte (0b1011) can store the same information.

## Problem:

If density is too low, bitmaps lose their edge. E.g., it is wasteful to represent {1,32000,64000} with a bitmap!

# *Bitmap compression*

- Keep the main benefits of bitmaps (e.g., bit-level parallelism)
- Support varied sets

Many solutions have  
been proposed!

Most popular  
solution:

run-length  
encoding (RLE)

Popularized by Oracle  
(BBC).

Many variations:  
WAH, EWAH,  
Concise.....

RLE-con  
used in n

acit  
BC).

RLE-compressed bitmaps are used in many production systems:

- Apache Hive,
- Oracle,
- Git...



Counting Objects

<http://githubengineering.com/counting-objects/>

# Main idea behind RLE:

0b0000....000000011111.....11111

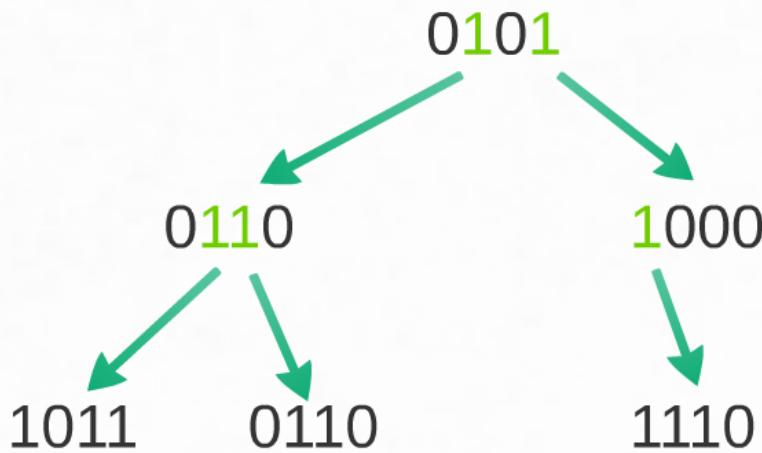
Then code it as

<number of zeros> <number of ones>

## Benefit of RLE:

- Good compression
- Pretty fast decoding (though limited by branch mispredictions?)

## Variation on RLE: hierarchical bitmaps



This is the compressed representation of a 4 x 4 x 4-bit bitmap containing many zeroes

SparseBitSet:  
<https://github.com/brettwooldridge/SparseBitSet>

Bitmagic:  
<http://bmagic.sourceforge.net/>

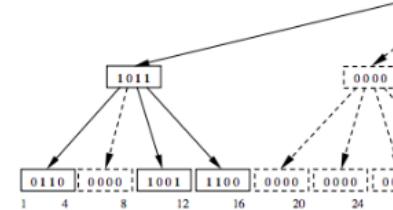


Fig. 1. A sing

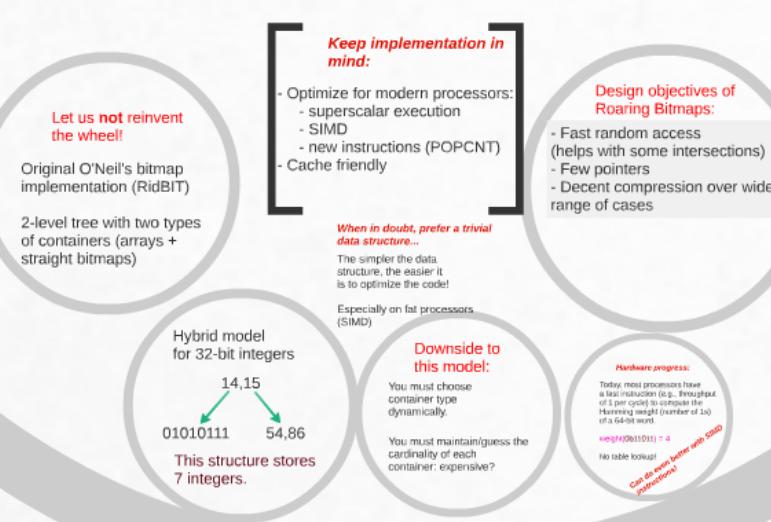
Downside of RLE : Difficult to do random access! Must start from the start always?

Downside of hierarchical : pointer overhead? Latency from tree traversal?

also common)

ways fast,  
intersection slower

alk...



## Design objectives of Roaring Bitmaps:

- Fast random access  
(helps with some intersections)
- Few pointers
- Decent compression over wide range of cases

on IN

essors:

CNT)

trivial

rs

## *Keep implementation in mind:*

- Optimize for modern processors:
  - superscalar execution
  - SIMD
  - new instructions (POPCNT)
- Cache friendly

***When in doubt, prefer a trivial  
data structure...***

The simpler the data  
structure, the easier it  
is to optimize the code!

Especially on fat processors  
(SIMD)



## *Array of containers...*

Decompose 32-bit space into 16-bit spaces (chunk). For each chunk, use best container:

- bitset (01001011)
- array ({1,20,144})
- runs ([0,10],[15,20])

Especially on fast processors  
(SIMD)

## Downside to this model:

You must choose  
container type  
dynamically.

You must maintain/guess the  
cardinality of each  
container: expensive?

## ***Hardware progress:***

Today, most processors have a fast instruction (e.g., throughput of 1 per cycle) to compute the Hamming weight (number of 1s) of a 64-bit word.

$$\text{weight}(0b11011) = 4$$

No table lookup!

*Can do even better with SIMD  
instructions!*

New hardware +  
well chosen algorithms  
  
make the hybrid model  
***fast.***

# ***Implementation challenge***

Instead of making everything fit in one model (bitset or array), we mix them.

Many optimization opportunities!

## Bitmap vs. Bitmap

**Intersection:** First compute cardinality of result, if low, construct directly an array (pain!), otherwise generate a bitmap (fast).

**Union:** Always generate a bitmap container (super fast)

## Array vs. Array

Intersection is always an array. If one array is much larger, we use  $O(n \log m)$  galloping.

**Union:** if some of cardinalities is large, guess that result should be bitmap. Do it (fast). If got it wrong, convert back to array (slow).

## ***Array vs. Bitmap***

Intersection is always array. Very fast:  
iterate over array and check bits in bitmap.

Union is always a bitmap: just set corresponding  
bits from array in bitmap. Super fast.

***To sum it up...***

Array vs. Bitmap: **super fast** (also common)

Bitmap vs. Bitmap: union is always fast,  
unlucky intersection slower

Array vs. Array : can be slow

*Run containers : For another talk...*

## *Example* **(Census 1881)**

	Concise	EWAH	Roaring	WAH
size (bits/int)	2.9	3.3	<b>2.7</b>	2.9
intersections (relative timings)	460	150	<b>1.0</b>	370
unions (relative timings)	43	43	<b>1.0</b>	38

See [https://github.com/RoaringBitmap/  
RoaringBitmap/tree/master/jmh](https://github.com/RoaringBitmap/RoaringBitmap/tree/master/jmh)

Several ports: (available NOW)

Go (complete, supported by us)

[github.com/RoaringBitmap/roaring](https://github.com/RoaringBitmap/roaring)

C++, C, Python, Cython, Rust,  
Haskell, C#

See <http://roaringbitmap.org/>

## Future/ongoing work:

- Official C port (highly optimized!)
- C#/LINQ
- Support for 64-bit integers
- Parallelization
- GPUs
- More in-depth comparisons



# ***Conclusion***

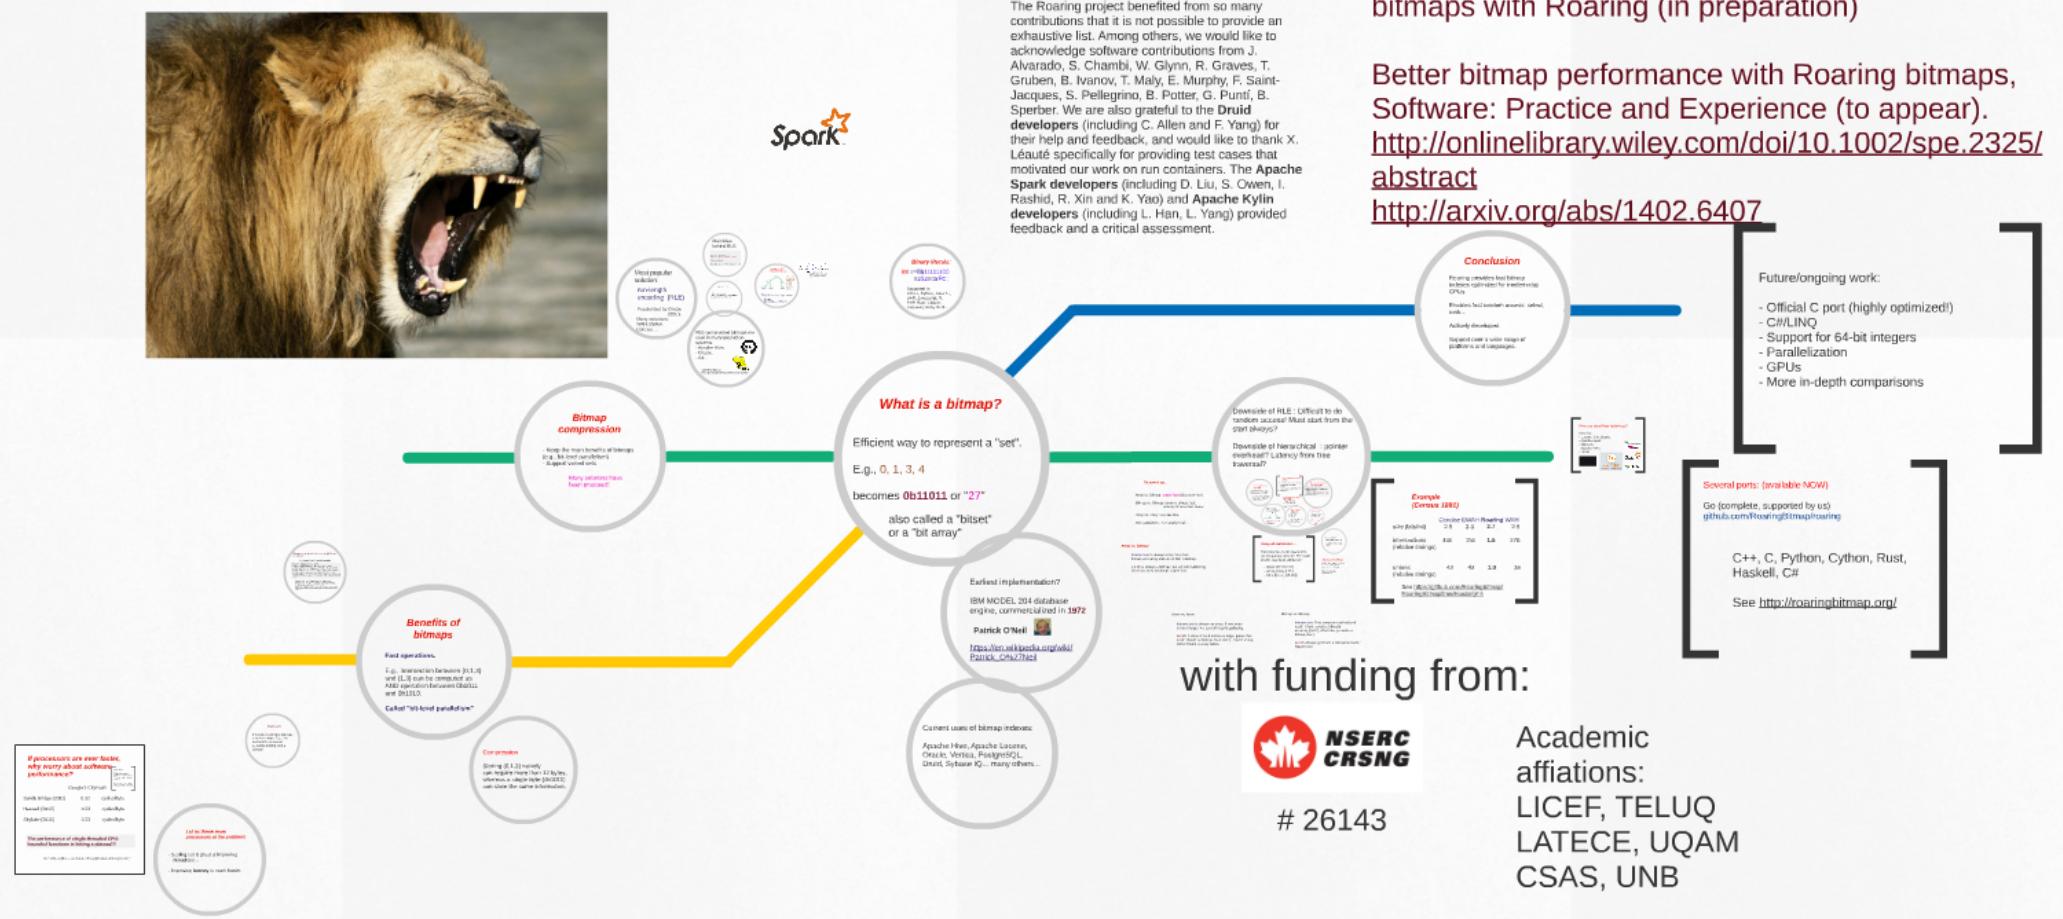
Roaring provides fast bitmap indexes optimized for modern-day CPUs.

Enables fast random access: select, rank...

Actively developed.

Support over a wide range of platforms and languages.

January 2016



<http://roaringbitmap.org/>  
<https://github.com/RoaringBitmap>

# Roaring Bitmaps

Consistently faster and smaller compressed bitmaps with Roaring (in preparation)

Better bitmap performance with Roaring bitmaps, Software: Practice and Experience (to appear).

<http://onlinelibrary.wiley.com/doi/10.1002/spe.2325/abstract>

<http://arxiv.org/abs/1402.6407>

with funding from:



# 26143

Academic affiliations:  
LICEF, TELUQ  
LATECE, UQAM  
CSAS, UNB