

# Unicode at gigabytes per second

Daniel Lemire with Wojciech Muła and John Keiser  
professor, Université du Québec (TÉLUQ)  
Montreal 🇨🇦

blog: <https://lemire.me>

twitter: [@lemire](https://twitter.com/lemire)

GitHub: <https://github.com/lemire/>

credit for figures: Wojciech Muła

! many other contributors!

# From characters to bits

Morse code

- A : 0 1
- B : 1 0 0 0
- C : 1 0 1 0

26 letters.

# Fixed-length codes

- Baudot code (~1860). 5 bits.
- Hollerith code (~1896). 6 bits.
- American Standard-Code for Information Interchange or ASCII (~1961). 7 bits. 128 characters.

| Dec | Hex | Chr | Dec | Hex | Chr   | Dec | Hex | Chr | Dec | Hex | Chr |
|-----|-----|-----|-----|-----|-------|-----|-----|-----|-----|-----|-----|
| 0   | 00  | NUL | 32  | 20  | Space | 64  | 40  | @   | 96  | 60  | `   |
| 1   | 01  | SOH | 33  | 21  | !     | 65  | 41  | A   | 97  | 61  | a   |
| 2   | 02  | STX | 34  | 22  | "     | 66  | 42  | B   | 98  | 62  | b   |
| 3   | 03  | ETX | 35  | 23  | #     | 67  | 43  | C   | 99  | 63  | c   |
| 4   | 04  | EOT | 36  | 24  | \$    | 68  | 44  | D   | 100 | 64  | d   |
| 5   | 05  | ENQ | 37  | 25  | %     | 69  | 45  | E   | 101 | 65  | e   |
| 6   | 06  | ACK | 38  | 26  | &     | 70  | 46  | F   | 102 | 66  | f   |
| 7   | 07  | BEL | 39  | 27  | '     | 71  | 47  | G   | 103 | 67  | g   |
| 8   | 08  | BS  | 40  | 28  | (     | 72  | 48  | H   | 104 | 68  | h   |
| 9   | 09  | HT  | 41  | 29  | )     | 73  | 49  | I   | 105 | 69  | i   |
| 10  | 0A  | LF  | 42  | 2A  | *     | 74  | 4A  | J   | 106 | 6A  | j   |
| 11  | 0B  | VT  | 43  | 2B  | +     | 75  | 4B  | K   | 107 | 6B  | k   |
| 12  | 0C  | FF  | 44  | 2C  | ,     | 76  | 4C  | L   | 108 | 6C  | l   |
| 13  | 0D  | CR  | 45  | 2D  | -     | 77  | 4D  | M   | 109 | 6D  | m   |
| 14  | 0E  | SO  | 46  | 2E  | .     | 78  | 4E  | N   | 110 | 6E  | n   |
| 15  | 0F  | SI  | 47  | 2F  | /     | 79  | 4F  | O   | 111 | 6F  | o   |
| 16  | 10  | DLE | 48  | 30  | 0     | 80  | 50  | P   | 112 | 70  | p   |
| 17  | 11  | DC1 | 49  | 31  | 1     | 81  | 51  | Q   | 113 | 71  | q   |
| 18  | 12  | DC2 | 50  | 32  | 2     | 82  | 52  | R   | 114 | 72  | r   |
| 19  | 13  | DC3 | 51  | 33  | 3     | 83  | 53  | S   | 115 | 73  | s   |
| 20  | 14  | DC4 | 52  | 34  | 4     | 84  | 54  | T   | 116 | 74  | t   |
| 21  | 15  | NAK | 53  | 35  | 5     | 85  | 55  | U   | 117 | 75  | u   |
| 22  | 16  | SYN | 54  | 36  | 6     | 86  | 56  | V   | 118 | 76  | v   |
| 23  | 17  | ETB | 55  | 37  | 7     | 87  | 57  | W   | 119 | 77  | w   |
| 24  | 18  | CAN | 56  | 38  | 8     | 88  | 58  | X   | 120 | 78  | x   |
| 25  | 19  | EM  | 57  | 39  | 9     | 89  | 59  | Y   | 121 | 79  | y   |
| 26  | 1A  | SUB | 58  | 3A  | :     | 90  | 5A  | Z   | 122 | 7A  | z   |
| 27  | 1B  | ESC | 59  | 3B  | ;     | 91  | 5B  | [   | 123 | 7B  | {   |
| 28  | 1C  | FS  | 60  | 3C  | <     | 92  | 5C  | \   | 124 | 7C  |     |
| 29  | 1D  | GS  | 61  | 3D  | =     | 93  | 5D  | ]   | 125 | 7D  | }   |
| 30  | 1E  | RS  | 62  | 3E  | >     | 94  | 5E  | ^   | 126 | 7E  | ~   |
| 31  | 1F  | US  | 63  | 3F  | ?     | 95  | 5F  | _   | 127 | 7F  | DEL |

# Too many fixed-length codes!

- IBM: Binary Coded Decimal Interchange Code. 6 bits.
- IBM: Extended Binary Coded Decimal Interchange Code or EBCDIC. 8 bits.
- ISO 8859 (~1987). 8 bits. European.
- Thai (TIS 620), Indian languages (ISCII), Vietnamese (VISCII) and Japanese (JIS X 0201).
- Windows character sets, Mac character sets.

# Unicode (late 1980s)

- Extends ASCII.
- Universal.
- Replaces all other standards.
- Typography, full localisation, extensible.

# Unicode: how many bits?

- 16 bits ought to be enough?
- Numerical range from 0x000000 to 0x10FFFF.
- Would need 20 to 21 bits.

# UTF-16 and UTF-8

Two main formats.

UTF-16: Java, C#, Windows

UTF-8: XML, JSON, HTML, Go, Rust, Swift



# UTF-16 and UTF-8

| character range                | UTF-8 bytes | UTF-16 bytes |
|--------------------------------|-------------|--------------|
| ASCII (0000-007F)              | 1           | 2            |
| latin (0080-07FF)              | 2           | 2            |
| asiatic (0800-D7FF, E000-FFFF) | 3           | 2            |
| supplemental (010000-10FFFF)   | 4           | 4            |

# UTF-16

- 16-bit words.
- characters in 0000-D7FF and E000-FFFF, stored as 16-bit values---using two bytes.
- characters in 010000-10FFFF are stored using a 'surrogate pair'.
- Comes in two flavours (little and big endian at the 16-bit level).

## UTF-16 (surrogate pair)

- first word in D800-DBFF.
- second word in DC00-DFFF.
- character value is 10 least significant bits of each---second element is least significant.
- add 0x10000 to the result.

# UTF-8

- 8-bit words (no endianness)
- One 'leading' byte followed by 0 to 3 bytes.


# UTF-8 format

- Most significant bit of leading is zero, ASCII: [01000001].
- 3 most significant bits 110, two-byte sequence: [11000100] [10000101].
- 4 most significant bits 1100, three-byte sequence.
- 5 most significant bits 11000, four-byte sequence.
- Non-leading bytes have 10 as the two most significant bits.


# UTF-8 validation rules

- The five most significant bits of any byte cannot be all ones.
- The leading byte must be followed by the right number of continuation bytes.
- A continuation byte must be preceded by a leading byte.
- The decoded character must be larger than 7F for two-byte sequences, larger than 7FF for three-byte sequences, and larger than FFFF for four-byte sequences.
- The decoded code-point value must be less than 110000
- The code-point value must not be in the range D800-DFFF.

# UTF-8/UTF-16 comparison (ASCII)




# UTF-8/UTF-16 comparison (2-bytes)






# UTF-8/UTF-16 comparison (3-bytes)



# UTF-8/UTF-16 comparison (4-bytes)



## UTF-8/UTF-16 transcoding

- Must convert (transcode) from one format to the other format, while validating the input.

## Some numbers

- bandwidth between node instances: over 3 GB/s
- PCIe 4.0 disks (and PlayStation 5): over 5 GB/s
- Popular C++ transcoding library (ICU): ~1 GB/s

# Gigabytes per second?

- x64, ARM, POWER: have SIMD instructions.

|                        | <b>UTF-8 to UTF-16</b> | <b>UTF-16 to UTF-8</b> | <b>validation</b> | <b>table size</b> |
|------------------------|------------------------|------------------------|-------------------|-------------------|
| Cameron's u8u16 (2008) | yes                    | no                     | yes               | N/A               |
| Inoue et al. (2008)    | partial                | no                     | no                | 105 kB            |
| simdutf                | yes                    | yes                    | yes               | 20 kB             |

Software implementations (no formal paper): Goffart (2012) and Gatilov (2019)

# Vectorized permutation

- Can permute blocks of 16 bytes (or 32 bytes) using a single cheap instruction.
- Need a precomputed shuffle mask.
- data : [a b c d e f g]
- shuffle mask : [3 1 0 3 3 2 -1] (indexes)
- result : [d b a d d c 0]
- Conversely may be used as a form of vectorized table lookup.

## UTF-8 to UTF-16 transcoding (core)

- Take a block of bytes.
- Continuation bytes (leading bits 10, less than -64)
- Non-continuation bytes are leading bytes
- Bytes before a leading byte end a character
- Build a bitmap
- Use the bitmap in a lookup table



# UTF-8 to UTF-16 transcoding (example)

Start with...

[01000001] ([11000100] [10000101])

[01100011] ([11000011] [10000011]) [01101100] ([11000101] [10111010])

We have 9 bytes. Build a 9-bit bitmap where '1' means the end of a character

101101101

Use this as index in a table.

## UTF-8 to UTF-16 transcoding (table)


- If using 12-byte blocks, need 4096-long table.
- Each entry points to a shuffle mask and number of consumed bytes.

# UTF-8 to UTF-16 transcoding (cases)

Shuffle masks are sorted into 'cases'.

1. First 64 cases correspond to 1-byte or 2-byte characters only.
2. Next 81 cases correspond to 1, 2 or 3 bytes per character.
3. Next 64 cases correspond to general case (1 to 4 bytes).

Each case corresponds to a code path.



## UTF-8 to UTF-16 transcoding (more tricks)

1. Load blocks of 64 bytes.
2. Check for fast paths (e.g. all ASCII).
3. Eat 12 bytes at a time within 64 bytes.
4. Add a few fast path (e.g., all ASCII, all 2-byte, all 3-byte).

# UTF-8 to UTF-16 transcoding (validation)

Given a 64-byte block, we can use a fast vectorized validation routine.

- Validating UTF-8 In Less Than One Instruction Per Byte, Software: Practice and Experience 51 (5), 2021

## UTF-8 to UTF-16 transcoding (core algo)

- You can identify most UTF-8 errors by looking at sequences of 3 nibbles (4-bit).
- E.g., ASCII followed by continuation, leading not followed by continuation byte.

Do three lookups (using shuffle mask) and compute a bitwise AND. We call this vectorized classification.

# Simplified vectorized classification

- Suppose you want to find all instances where value 3 is followed by value 1 or 2.
- Create two lookup tables.
- One for first nibble [0,0,0,1,0,0,0,0,0,0,0,0,0,0,0]
- second nibble [0,1,1,0,0,0,0,0,0,0,0,0,0,0,0]
- Lookup first nibble in table, lookup second, compute bitwise AND.
- If result is 1, you have a match.
- Can do this in parallel over many values.



# Fancier vectorized classification

- Suppose you want to find all instances where value 3 is followed by value 1 or 2. Value 5 followed by 0. Value 6 followed by 10.
- Create two lookup table2.
- One for first nibble [0,0,0,1,0,2,4,0,0,0,0,0,0,0,0,0]
- second nibble [2,1,1,0,0,0,0,0,0,0,4,0,0,0,0,0]
- Lookup first nibble in table, lookup second, compute bitwise AND.

Array of nibbles:

- original: [a0 a1 a2 a3 a4 ...]
- shift: [a1 a2 a3 a4 ...]
- shift: [a2 a3 a4 ...]
- $f([a0\ a1\ a2\ a3\ a4\ \dots]) \text{ AND } g([a1\ a2\ a3\ a4\ \dots]) \text{ AND } g([a2\ a3\ a4\ \dots])$

## UTF-16 to UTF-8

The other direction (from UTF-16 to UTF-8) is somewhat easier!

## UTF-16 to UTF-8 (ASCII)

If all 16-bit words are ASCII (0000-007F), use a fast routine: 16 bytes into 8 'packed' bytes.

## UTF-16 to UTF-8 (0000-07FF)

If all 16-bit words are in (0000-07FF)... build an 8-bit bitset indicating which 16-byte words are ASCII (0000-007F), load a shuffle mask, permute and patch.

## UTF-16 to UTF-8 (0000-07FF, E000-FFFF)

If all 16-bit words are in the ranges 0000-D7FF, E000-FFFF, we use another similar specialized routine to produce sequences of one-byte, two-byte and three-byte UTF-8 characters.

Otherwise, when we detect that the input register contains at least one part of a surrogate pair, we fall back to a conventional/scalar code path.


# Experiments

- AMD processor (AMD EPYC 7262, Zen 2 microarchitecture, 3.39 GHz) and GCC10.
- International Components for Unicode (UCI)
- u8u16 library
- lipsum text in various languages


# ASCII transcoding

|         | UTF-8 to UTF-16 | UTF-16 to UTF-8 |
|---------|-----------------|-----------------|
| simdutf | 20 GB/s         | 36 GB/s         |
| UCI     | 1 GB/s          | 2 GB/s          |





(a) UTF-8 to UTF-16 transcoding



(b) UTF-16 to UTF-8 transcoding

# Software

<https://github.com/simdutf/simdutf>

- Open source, no patent.
- ARM NEON, SSE, AVX...
- Support runtime dispatch: adapts to your CPU.
- Easy to use: drop `simdutf.cpp` and `simdutf.h` in your project.
- Compiles to tens of kilobytes.

## Further reading

- Lemire, Daniel and Wojciech Muła , Transcoding Billions of Unicode Characters per Second with SIMD Instructions, Software: Practice and Experience (to appear) <https://r-libre.telug.ca/2400/>
- Blog: <https://lemire.me/blog/>