

# Binary Fuse Filters: Fast and Tiny Immutable Filters

Daniel Lemire

professor, Data Science Research Center

Université du Québec (TÉLUQ)

Montreal

blog: <https://lemire.me>

twitter: [@lemire](https://twitter.com/lemire)

GitHub: <https://github.com/lemire/>

# Probabilistic filters?

- Is  $x$  in the set  $S$ ?
- Maybe or *definitively not*

# Usage scenario?

- We have this expensive *database*. Querying it cost you.
- Most queries should not end up in the data.
- We want a small 'filter' that can prune out queries.

# Theoretical bound

- Given  $N$  elements in the set
- Spend  $k$  bits per element
- Get a false positive rate of  $1/2^k$

# Usual constraints

- Fixed initial capacity
- Difficult to update safely without access to the set
- To get a 1% false-positive rate:  $\approx 8$  bits?

# Hash function

- From any objet in the *universe* to a *word* (e.g., 64-bit word)
- Result looks random

```
uint64_t murmur64(uint64_t h) {  
    h ^= h >> 33;  
    h *= UINT64_C(0xff51afd7ed558ccd);  
    h ^= h >> 33;  
    h *= UINT64_C(0xc4ceb9fe1a85ec53);  
    h ^= h >> 33;  
    return h;  
}
```

# Conventional Bloom filter

- Start with a bitset  $B$ .
- Using  $k$  hash functions  $f_1, f_2, \dots$



# Adding an element

- Given an object  $x$  from the set, set up to  $k$  bits to 1
- $B[f_1(x)] \leftarrow 1, B[f_2(x)] \leftarrow 1, \dots$

# Checking an element

- Given an object  $x$  from the universe, set up to  $k$  bits to 1
- $(B[f_1(x)] = 1) \text{ AND } (B[f_2(x)] = 1) \text{ AND } \dots$

# Checking an element: implementation

- Typical implementation is *branchy*
- If not  $(B[f_1(x)] = 1)$ , return false
- If not  $(B[f_2(x)] = 1)$ , return false
- ...
- return true


```
uint64_t hash = hasher(key);
uint64_t a = (hash >> 32) | (hash << 32);
uint64_t b = hash;
for (int i = 0; i < k; i++) {
    if ((data[reduce(a, length)] & getBit(a)) == 0) {
        return NotFound;
    }
    a += b;
}
return Found;
```


# False positive rate

| bits per element | hash functions | fpp   |
|------------------|----------------|-------|
| 9                | 6              | 1.3%  |
| 10               | 7              | 0.8%  |
| 12               | 8              | 0.3%  |
| 13               | 9              | 0.2%  |
| 15               | 10             | 0.07% |
| 16               | 11             | 0.04% |

## Bloom filters: upsides

- Fast construction
- Flexible: excess capacity translates into lower false positive rate
- Degrades smoothly to a useless but 'correct' filter








## Bloom filters: downsides

- 44% above the theoretical minimum in storage
- Slower than alternatives (lots of memory accesses)



# Memory accesses

| number of hash functions | cache misses (miss) | cache misses (hit) |
|--------------------------|---------------------|--------------------|
| 8                        | 3.5                 | 7.5                |
| 11                       | 3.8                 | 10.5               |

(Intel Ice Lake processor, out-of-cache filter)

# Mispredicted branches

| number of hash functions | all out | all in |
|--------------------------|---------|--------|
| 8                        | 0.95    | 0.0    |
| 11                       | 0.95    | 0.0    |

(Intel Ice Lake processor, out-of-cache filter)

# Performance

| number of hash functions | always out (cycles/entry) | always in (cycles/entry) |
|--------------------------|---------------------------|--------------------------|
| 8                        | 135                       | 170                      |
| 11                       | 140                       | 230                      |

(Intel Ice Lake processor, out-of-cache filter)

# Blocked Bloom filters

- Same as a Bloom filters, but for a given object, put all bits in one cache line
- Optional: Use SIMD instructions to reduce instruction count

## Blocked Bloom filters: pros/cons

- Stupidly fast in both construction and queries
- ~56% above the theoretical minimum in storage

```
auto hash = hasher_(key);  
uint32_t bucket_idx = reduce(rotl64(hash, 32), bucketCount);  
__m256i mask = MakeMask(hash);  
__m256i bucket = directory[bucket_idx];  
return _mm256_testc_si256(bucket, mask);
```



# Binary fuse filters

- Based on theoretical work by Dietzfelbinger and Walzer
- Immutable datastructure: build it once
- Fill it to capacity
- Fast construction
- Fast and simple queries

## Arity : 3-wise, 4-wise

- 3-wise version has three hits, 12% overhead
- 4-wise version has four hits, 8% overhead

# Queries are silly

- Have an array of *fingerprints* (e.g., 8-bit words)
- Compute 3 (or 4) hash functions:  $f_1(x)$ ,  $f_2(x)$ ,  $f_3(x)$
- Compute fingerprint function ( $f(x) \rightarrow$  8-bit word)
- Compute XOR and compare with fingerprint:

$$B[f_1(x)] \text{ XOR } B[f_2(x)] \text{ XOR } B[f_3(x)] = f(x)$$


```
bool contain(uint64_t key, const binary_fuse_t *filter) {  
    uint64_t hash = mix_split(key, filter->Seed);  
    uint8_t f = fingerprint(hash);  
    binary_hashes_t hashes = hash_batch(hash, filter);  
    f ^= filter->Fingerprints[hashes.h0] ^ filter->Fingerprints[hashes.h1] ^  
        filter->Fingerprints[hashes.h2];  
    return f == 0;  
}
```

|                    | <b>cache misses</b> | <b>mispredictions</b> |
|--------------------|---------------------|-----------------------|
| 3-wise binary fuse | 2.8                 | 0.0                   |
| 4-wise binary fuse | 3.7                 | 0.0                   |


(Intel Ice Lake processor, out-of-cache filter)

|                     | <b>always out<br/>(cycles/entry)</b> | <b>always in<br/>(cycles/entry)</b> | <b>bits per<br/>entry</b> |
|---------------------|--------------------------------------|-------------------------------------|---------------------------|
| Bloom $k = 8$       | 135                                  | 170                                 | 12                        |
| 3-wise bin.<br>fuse | 85                                   | 85                                  | 9.0                       |
| 4-wise bin.<br>fuse | 100                                  | 100                                 | 8.6                       |

(Intel Ice Lake processor, out-of-cache filter)



(a) Relative space usage



(b) Query time

# Construction 1

- Start with array for fingerprints containing slightly more fingerprints than you have elements in the set
- Divide the array into segments (e.g., 300 disjoint)
- Number of fingerprints in segment: power of two (hence *binary*)



## Construction 2

- Map each object  $x$  in set, to locations  $B[f_1(x)]$ ,  $B[f_2(x)]$ ,  $B[f_3(x)]$
- The locations should be in three consecutive segments (so relatively nearby in memory).

## Construction 3

- At the end, each location  $B[i]$  is associated with some number of objects from the set

## Construction 4

- Find a location mapped from a single set element  $x$ , e.g.,  $B[f_1(x)]$
- Record this location which is owned by  $x$
- Remove the mapping of  $x$  to locations  $B[f_1(x)]$ ,  $B[f_2(x)]$ ,  $B[f_3(x)]$
- Repeat


## Construction 5

- Almost always, the construction terminates after one trial
- Go through the matched keys, in reverse order, and set (e.g.)

$$B[f_1(x)] = f(x) \text{ XOR } B[f_2(x)] \text{ XOR } B[f_3(x)]$$

# Construction: Performance

- Implemented naively: terrible performance (random access!!!)
- Before the construction begins, sort the elements of the sets according to the segments they are mapped to.
- This greatly accelerates the construction



# How does the performance scale with size?

For warm small filters, number of access is less important.  
Becomes more computational.

For large cold filters, accesses are costly.

# 10M entries

|                      | ns/query (all out) | ns/query (all in) | fpp   | bits per entry |
|----------------------|--------------------|-------------------|-------|----------------|
| Bloom                | 17                 | 14                | 0.32% | 12.0           |
| Blocked Bloom (NEON) | 3.8                | 3.8               | 0.6%  | 12.8           |
| 3-wise bin. fuse     | 3.5                | 3.5               | 0.39% | 9.0            |
| 4-wise bin. fuse     | 4.0                | 4.0               | 0.39% | 8.6            |

(Apple M2)



# 100M entries

|                      | ns/query (all out) | ns/query (all in) | fpp   | bits per entry |
|----------------------|--------------------|-------------------|-------|----------------|
| Bloom                | 38                 | 33                | 0.32% | 12.0           |
| Blocked Bloom (NEON) | 11                 | 11                | 0.6%  | 12.8           |
| 4-wise bin. fuse     | 17                 | 17                | 0.39% | 9.0            |
| 4-wise bin. fuse     | 20                 | 20                | 0.39% | 8.6            |

(Apple M2)

# Compressibility (zstd)

|                  | bits per entry (raw) | bits per entry (zstd) |
|------------------|----------------------|-----------------------|
| Bloom $k = 8$    | 12.0                 | 12.0                  |
| 3-wise bin. fuse | 9.0                  | 8.59                  |
| 4-wise bin. fuse | 8.60                 | 8.39                  |
| theory           | 8.0                  | 8.0                   |

# **Sending compressed filters**

Compressed (zstd) binary fuse filters can be within 5% of the theoretical minimum.

## Some links

- Bloom filters in Go: <https://github.com/bits-and-blooms/bloom>
- Binary fuse filters in Go: <https://github.com/FastFilter/xorfilter>
- Binary fuse filters in C: [https://github.com/FastFilter/xor\\_singleheader](https://github.com/FastFilter/xor_singleheader)
- Binary fuse filters in Java: [https://github.com/FastFilter/fastfilter\\_java](https://github.com/FastFilter/fastfilter_java)
- Giant benchmarking platform: [https://github.com/FastFilter/fastfilter\\_cpp](https://github.com/FastFilter/fastfilter_cpp)

## Other Links

- Blog <https://lemire.me/blog/>
- Twitter: @lemire
- GitHub: <https://github.com/lemire>