

# Accurate and efficient software microbenchmarks

Daniel Lemire

professor, Data Science Research Center

Université du Québec (TÉLUQ)

Montreal 

blog: <https://lemire.me>

twitter: [@lemire](https://twitter.com/lemire)

GitHub: <https://github.com/lemire/>

# Background

- *Fastest JSON parser in the world* (on commodity processors):  
<https://github.com/simdjson/simdjson>
- First to parse JSON files at gigabytes per second

# Where is the code ?

All code for this talk is online (reproducible!!!)

<https://github.com/lemire/talks/tree/master/2023/performance/code>

# How fast is your disk?

PCIe 4 drives: 5 GB/s reading speed (sequential)

PCIe 5 drives: 10 GB/s reading speed (sequential)



# CPU Frequencies are stagnating

| architecture   | availability | max. frequency |
|----------------|--------------|----------------|
| Intel Skylake  | 2015         | 4.5 GHz        |
| Intel Ice Lake | 2019         | 4.1 GHz        |

# Fact

Single-core processes are often CPU bound

# Solution?

Optimize the software.

Incremental optimization, how do you know that you are on the right track?

# Hypothesis

This software change (commit) improves our performance.

# Simple


Measure time elapsed before, time elapsed after.

# Complex system

Software systems are complex systems: changes can have unexpected consequences.

# JIT

## *Virtual Machine Warmup Blows Hot and Cold*



# System calls

System calls (especially IO) may dominate, assume that they remain constant. Idem with multicore and multi-system processes.

# Data access


data structure layout changes can trigger expensive loads, assume that we keep that constant.

# Tiny functions

Uncertainty principle: by measuring you are affecting the execution so that you cannot measure safely tiny functions.

# Take statically compiled code

Transcoding UTF-16 to UTF-8 of an 80kB Arabic string using the simdutf library (NEON kernel).



# Use the average?

Let  $t$  be the true value and let  $\epsilon$  be the noise distribution (variance  $\sigma^2$ ).

We seek  $t$ .

# Repeated measures increase accuracy


Measures are  $t + \epsilon_1, t + \epsilon_2, t + \epsilon_3, \dots$

Sum is  $Nt + (\sum_i \epsilon_i)$ . Variance is  $N\sigma^2$ .

Average is  $t + (\sum_i \epsilon_i)/N$ . Variance is  $\sigma^2/N$ . Standard deviation of  $\frac{\sigma}{\sqrt{N}}$ .

# Simulation

```
mu, sigma = 10000, 5000
for N in range(20, 2000+1):
    s = [sum(np.random.default_rng().normal(mu, sigma, N))/N for i in range(30)]
    print(N, np.std(s))
```




# Actual measurements


```
// returns the average
double transcode(const std::string& source, size_t iterations);

...

for(size_t i = iterations_start; i <= iterations_end; i+=step) {
    std::vector<double> averages;
    for(size_t j = 0; j < 30; j++) { averages.push_back(transcode(source, i)); }
    std::cout << i << "\t" << compute_std_dev(averages) << std::endl;
}
```



# Sigma events



- 1-sigma is 32%
- 2-sigma is 5%
- 3-sigma is 0.3% (once ever 300 trials)
- 4-sigma is 0.00669% (once every 15000 trials)
- 5-sigma is 5.9e-05% (once every 1,700,000 trials)
- 6-sigma is 2e-07% (once every 500,000,000)

$$e^{-n^2/2} / (n * \sqrt{\pi/2}) \times 100 \text{ for } n > 3$$


# Measuring sigma events

Take 300 measures after warmup, and measure the worst relative deviation


```
$ for i in {1..10}; do sudo ./sigma_test; done  
4.56151  
4.904  
7.43446  
5.73425  
9.89544  
12.975  
3.92584  
3.14633  
4.91766  
5.3699
```

# What if we dealt with log-normal distributions?

Log-Normal Distribution ( $\mu=1, \sigma=1$ )



```
for N in range(20, 2000+1):  
    s = [sum(np.random.default_rng().lognormal(1, 4, N))/N for i in range(30)]  
    print(N, np.std(s))
```



# What if we measured the minimum?

Relative standard deviation ( $\sigma/\mu$ )

| <b>N</b> | <b>average</b> | <b>minimum</b> |
|----------|----------------|----------------|
| 200      | 3.44%          | 1.38%          |
| 2000     | 2.66%          | 1.19%          |
| 10000    | 2.95%          | 1.27%          |

- The minimum is easier to measure to 1% accuracy.

# CPU performance counters

Processors have *zero-overhead* counters recording instruction retired, actual cycles, and so forth.

No need to freeze the CPU frequency: you can measure it.


# Limitations

- You can only measure so many things (2, 4 metrics, not 25)
- Required privileged access (e.g., root)

# Counters in the cloud

- x64: Requires at least a full CPU
- ARM Graviton: generally available but limited number (e.g., 2 counters)

# Instruction counts are accurate



# Using performance counters

- Java instruction counters: <https://github.com/jvm-profiling-tools/async-profiler>
- C/C++: instruction counters are available through the Linux kernel
- [Go instruction counters](#)

# Generally, fewer instructions means faster code

- Some instructions are more expensive than others (e.g., division).
- Data dependency can make instruction counts less relevant.
- Branching can artificially lower instruction count.

If you are adding speculative branching, make sure your test input is large.

```
while (howmany != 0) {  
    val = random();  
    if( val is an odd integer ) {  
        out[index] = val;  
        index += 1;  
    }  
    howmany--;  
}
```

## 2000 'random' elements, AMD Rome

| trial | mispredicted branches |
|-------|-----------------------|
| 1     | 50%                   |
| 2     | 18%                   |
| 3     | 6%                    |
| 4     | 2%                    |
| 5     | 1%                    |
| 6     | 0.3%                  |
| 7     | 0.15%                 |
| 8     | 0.15%                 |
| 9     | 0.1%                  |

## Take away 1

- Computational microbenchmarks can have log-normal distributions.
- Consider measuring the *minimum* instead of the *average*.

## Take away 2

- Benchmarking often is good
- Long-running benchmarks are not necessarily more accurate.
- Prefer cheap, well-designed benchmarks.

# Links

- Blog <https://lemire.me/blog/>
- Twitter: @lemire
- GitHub: <https://github.com/lemire>