


## Algorithms for Modern Processor Architectures

Daniel Lemire, professor  
Université du Québec (TÉLUQ)  
Montréal 🇨🇦

blog: <https://lemire.me>

X: [@lemire](#)

GitHub: <https://github.com/lemire/>



## Disk at gigabytes per second


**Micron shows off world's fastest PCIe 6.0 SSD, hitting 27 GB/s speeds — Astera Labs PCIe 6.0 switch enables impressive sequential reads**

News

By [Sunny Grimm](#) published March 8, 2025


The next-gen of networking and storage is hitting the trade shows

## PCI Express Bandwidth Comparison



# High Bandwidth Memory


- Xeon Max processors contain 64 GB of HBM
- Bandwidth 800 GB/s



|           |           |           |           |           |           |           |           |
|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia |
| Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia |
| Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia |
| Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia | Wikipedia |

# Some numbers


- Time is discrete: clock cycle
- Processors: 4 GHz ( $4 \times 10^9$  cycles per second)
- One cycle is 0.25 nanoseconds
- light: 7.5 centimeters per cycle
- One byte per cycle: 4 GB/s
- **Easily CPU bound**



# Frequencies and transistors

| processor     | year | frequency | transistors    |
|---------------|------|-----------|----------------|
| Pentium 4     | 2000 | 3.8 GHz   | 0.040 billions |
| Intel Haswell | 2013 | 4.4 GHz   | 1.4 billions   |
| Apple M1      | 2020 | 3.2 GHz   | 16 billions    |
| Apple M2      | 2022 | 3.49 GHz  | 20 billions    |
| Apple M3      | 2024 | 4.05 GHz  | 25 billions    |
| Apple M4      | 2024 | 4.5 GHz   | 28 billions    |
| AMD Zen 5     | 2024 | 5.7 GHz   | 50 billions    |

Year vs number of transistors




# Moore's Law: The number of transistors on microchips has doubled every two years



Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

## Transistor count



# Where do the transistors go?

- More cores
- More superscalar execution (more instructions per cycle)
- Better speculative execution (→ more instructions per cycle)
- More cache, more memory-level parallelism (→ more instructions per cycle)
- Better data-level parallelism (SIMD) (→ fewer instructions)

# Superscalar execution

| processor       | year | arithmetic logic units | SIMD units     |
|-----------------|------|------------------------|----------------|
| Pentium 4       | 2000 | 2                      | $2 \times 128$ |
| AMD Zen 2       | 2019 | 4                      | $2 \times 256$ |
| Apple M*        | 2019 | 6+                     | $4 \times 128$ |
| Intel Lion Cove | 2024 | 6                      | $4 \times 256$ |
| AMD Zen 5       | 2024 | 6                      | $4 \times 512$ |


Moving to up to 4 load/store per cycle

# Parsing a number

- `1.3321321e-12` to `double`

```
double result;  
fast_float::from_chars(  
    input.data(), input.data() + input.size(), result);
```

# Parsing a number



# Lemire's Rule 1

Modern processors execute *nearly* as many instructions per cycle as you can supply.

- with caveats: branching, memory, and input/output

# Lemire's Corollary 1

In computational workloads (batches), minimizing instruction count is critical for achieving optimal performance.

## Lemire's Tips

1. Batch your work in larger units to save instructions.
2. Simplify the processing down to as few instructions as possible.
3. Every optimization counts (let go of 'hot spot' engineering).

# Going back to number parsing

- Our number parser: major browsers (Safari, Chrome), GCC (12+), C#, Rust, MySQL, Go
- About  $4\times$  faster than the conventional alternatives.
- How did we do it?

We massively reduced the number of CPU instructions required.

| function   | instructions  |
|------------|---------------|
| strtod     | $> 1000$      |
| our parser | $\approx 200$ |

*Reference:*

Number Parsing at a Gigabyte per Second, Software: Practice and Experience 51 (8), 2021

[https://github.com/fastfloat/fast\\_float](https://github.com/fastfloat/fast_float)

# SWAR

- Stands for SIMD within a register
- Use normal instructions, portable (in C, C++,...)
- A 64-bit register can be viewed as 8 bytes
- Requires some cleverness

## Check whether we have a digit


In ASCII/UTF-8, the digits 0, 1, ..., 9 have values 0x30, 0x31, ..., 0x39.

To recognize a digit:

- The high nibble should be 3.
- The high nibble should remain 3 if we add 6 (0x39 + 0x6 is 0x3f)

## ASCII processing flow

ASCII numbers



ASCII codes (hex)



Zero lower 4 bits (AND 0xF0)



Add 6, shift right 4



OR result (hex)



```
// load 8 input bytes into val
bool is_made_of_eight_digits_fast(uint64_t val) noexcept {
    return !(((val + 0x4646464646464646)
              | (val - 0x3030303030303030))
            & 0x8080808080808080));
}
```

to

```
add    rax, rdi
add    rdi, rdx
or     rax, rdi
test   rax, rdx
```

# Batching (unrolling)

6 to 7 instructions per multiplication

```
for (size_t i = 0; i < length; i++)  
    sum += x[i] * y[i];
```

3 to 5 instructions per multiplication

```
for (; i < length - 3; i += 4)  
    sum += x[i] * y[i]  
        + x[i + 1] * y[i + 1]  
        + x[i + 2] * y[i + 2]  
        + x[i + 3] * y[i + 3];
```

## Apple M4

Instructions



Regular: 6



Unrolled: 3.5

Cycles



Regular: 1.2



Unrolled: 1.0

## Intel Ice Lake

Instructions



Regular: 7




Unrolled: 5

Cycles



Regular: 1.6



Unrolled: 1.1


# Knuth's random shuffle

```
PROCEDURE shuffle(array)
  FOR j FROM |array| - 1 DOWN TO 1
    k ← random_integer(0, j)
    SWAP array[j] WITH array[k]
  END FOR
END PROCEDURE
```

# Batched random shuffle

- Draw one random number
- Compute two indices (with high probability)
- Reduces the instruction count
- Reduces the number of branches

## Results (Apple M4): Use a large array (8 MB).



*Reference:* Batched Ranged Random Integer Generation, Software: Practice and Experience 55 (1), 2025

[https://github.com/lemire/cpp\\_batched\\_random](https://github.com/lemire/cpp_batched_random)

# Branching

Hard-to-predict branches can derail performance

# Unicode (UTF-16)

- Code points from U+0000 to U+FFFF, a single 16-bit value.
- Beyond: a surrogate pair [U+D800 to U+DBFF] followed by U+DC00 to U+DFFF


# Validate

- Check whether we have a lone code unit ( $x \leq 0xD7FF \vee x \geq 0xDBFF$ ), if so ok
- Check whether we have the first part of the surrogate ( $0xD800 \leq x \leq 0xDBFF$ ) and if so check that we have the second part of a surrogate

# Validate


```
PROCEDURE validate_utf16(code_units)
  i ← 0
  WHILE i < |code_units|
    unit ← code_units[i]
    IF unit ≤ 0xD7FF OR unit ≥ 0xE000 THEN
      INCREMENT i
      CONTINUE
    IF unit ≥ 0xD800 AND unit ≤ 0xDBFF THEN
      IF i + 1 ≥ |code_units| THEN
        RETURN false
      next_unit ← code_units[i + 1]
      IF next_unit < 0xDC00 OR next_unit > 0xDFFF THEN
        RETURN false
      i ← i + 2 // Valid surrogate pair
      CONTINUE
    RETURN false
  RETURN true
```

# Performance results (Apple M4)



1 character per cycle might be just 4 GB/s (slower than disk)

# Performance results (Apple M4)




We are now barely at 1 GB/s!

# Speculative execution

- Processors *predict* branches
- They execute code *speculatively* (can be wrong!)

**How much can your processor learn?**



# Finite state machine to the rescue

- Can identify characters by the most significant 8 bits.
- Trivial finite state machine: default, has just encountered a high surrogate, or error.

```




static uint8_t transition_table[3][256] = {
    {...},
    {...},
    {...}
};

bool is_valid_utf16_ff(std::span<uint16_t> code_units) {
    uint8_t state = 0; // Start in Initial state
    for (auto code_unit : code_units) {
        uint8_t high_byte = code_unit >> 8;
        state = transition_table[state][high_byte];
    }
    return state == 0; // Valid only if we end in Initial state
}

```

# Performance results (Apple M4)


The finite-state approach can be  $7\times$  faster!



# Rules of thumb

1. Processors can 'learn' thousands of branches: benchmark over massive inputs.
2. Pick a solution without branches when it provides the same performance.

# Apple M4 can learn 10,000 random (0/1) branches.



# Pipelining

How does the processor manage to validate one UTF-16 character per cycle when it takes **many cycles** just to *load* the character?

| cycle | action        | action        | pizza en route |
|-------|---------------|---------------|----------------|
| 1     | order pizza A |               |                |
| 2     | order pizza B |               | A 🚚            |
| 3     | order pizza C |               | A 🚚, B 🚚       |
| 4     | order pizza D | eat pizza A 🍕 | B 🚚, C 🚚       |
| 5     | order pizza E | eat pizza B 🍕 | C 🚚, D 🚚       |
| 6     | order pizza F | eat pizza C 🍕 | D 🚚, E 🚚       |

# Little's Law


- Latency harms throughput
- Parallelism hides latency

$$\text{throughput} = \frac{\text{parallelism}}{\text{latency}}$$


# Memory-level parallelism


- Create large array of indices forming a cycle
- Start with [0,1,2,3,4]
- Shuffle so that no index can remain in place.
- Start with [4,0,3,2,1]
- Sandra Sattolo's algorithm
- This forms a random path

# 1 lane





## 2 lanes






# Consequence




# Bloom filter




# Bloom filter

8 hash functions, (Intel Ice Lake processor, out-of-cache filter)



Less than half the cache misses

# Bloom filter



# Data-level parallelism

# SIMD

- Stands for Single instruction, multiple data
- Allows us to process 16 bytes or more with one instruction
- Supported on all modern CPUs (phone, laptop)

# ASCII to lower case

```
For each character c
  If c - 'A' < 'Z' - 'A' then
    c = c + 'a' - 'A'
  EndIf
EndFor
```

# ASCII to lower case: 64 characters in 3 instructions

- Compute  $c - A$

```
__m512i ca = _mm512_sub_epi8(c, _mm512_set1_epi8('A'));
```

- Turn  $c - 'A' \leq Z - A$  into a mask

```
__mmask64 is_upper = _mm512_cmple_epu8_mask(ca, _mm512_set1_epi8('Z' - 'A'));
```

- Add  $a - A$  to  $c$  according to mask

```
__m512i to_lower = _mm512_mask_add_epi8(c, is_upper, c, to_lower)
```

# Deltas (C)




successive difference:

```
for (size_t i = 1; i < n; ++i) {  
    dst[i] = src[i] - src[i - 1];  
}
```




prefix sum:

```
for (size_t i = 1; i < n; ++i) {  
    dst[i] = dst[i - 1] + src[i];  
}
```

# Apple M4



# Now allow SIMD! (Autovectorization)






**Need to learn SIMD design magic !**

# UTF-16

- Write SIMD correction function (not just validation)
- Actually deployed in v8 (Google Chrome, Microsoft Edge)

# UTF-16, random (adversarial), Apple M4



- Most x64 processors have AVX2 (32-byte register)
- AMD Zen 5 has powerful AVX-512 (64-byte register)
- ARM has NEON + SVE/SVE2
- RISC-V has its vector instructions
- Loonson processes have AVX2-like instructions

# Interested? Check these projects

- simdjson: The fastest JSON parser in the world <https://simdjson.org>
  - Node.js, Electron,...
  - ClickHouse, WatermelonDB, Apache Doris, Meta Velox, Milvus, QuestDB, StarRocks
- simdutf: Unicode routines (UTF8, UTF16, UTF32) and Base64 <https://github.com/simdutf/simdutf>
  - Node.js, Bun, WebKit (Safari), Chromium (Chrome, Edge)


# Measurements

- We often assume that measurements (timings) are normally distributed.
- It is often an incorrect assumption.


# Measurements

- If your measurements are normally distributed, the 'error' falls off as  $1/\sqrt{N}$


Original normal distribution



Sum of 10 i.i.d. normals / 10





# Sigma events



- 1-sigma is 32%
- 2-sigma is 5%
- 3-sigma is 0.3% (once every 300 trials)
- 4-sigma is 0.00669% (once every 15000 trials)
- 5-sigma is 5.9e-05% (once every 1,700,000 trials)
- 6-sigma is 2e-07% (once every 500,000,000)
- $e^{-n^2/2} / (n * \sqrt{\pi/2}) \times 100$  for  $n > 3$

# What if we dealt with log-normal distributions?



# Real-world measurements

- You cannot assume normality
- Measurements are **not independent**.
- Reality: the absolute minimum is often a *reliable metric*
- Margin: difference between mean and minimum

# Conclusion

- Processors are getting much better! Wider!
- 'hot spot' engineering can fail, better to reduce overall instruction count.
- Branchy code can do well in synthetic benchmarks, but be careful.